

# Machine Learning Clustering and Dimensionality Reduction

## Assignment #1

Rhichard Koh

```
In [ ]: import pandas as pd
```

## Apply PCA to the train set

```
In [ ]: df = pd.read_csv('creditcard.csv')
df
```

```
Out[ ]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.0986
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.0851
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.2476
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.3774
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.2705
...	...	...	...	...	...	...	...	...	...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.3053
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.2948
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.7084
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.6791
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.4146

284807 rows × 31 columns

```
In [ ]: df.columns
```

```
Out[ ]: Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10',
        'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20',
        'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount',
        'Class'],
        dtype='object')
```

```
In [ ]: X = df.drop(columns=['Class'])
X
```

Out[ ]:

	Time	V1	V2	V3	V4	V5	V6	V7	
<b>0</b>	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.0986
<b>1</b>	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.0851
<b>2</b>	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.2476
<b>3</b>	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.3774
<b>4</b>	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.2705
...	...	...	...	...	...	...	...	...	...
<b>284802</b>	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.3053
<b>284803</b>	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.2948
<b>284804</b>	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.7084
<b>284805</b>	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.6791
<b>284806</b>	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.4146

284807 rows × 30 columns

In [ ]: `X.columns`

Out[ ]: Index(['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10', 'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20', 'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount'], dtype='object')

In [ ]: `y = df['Class']`  
y

Out[ ]:

0	0
1	0
2	0
3	0
4	0
..	
284802	0
284803	0
284804	0
284805	0
284806	0

Name: Class, Length: 284807, dtype: int64

In [ ]: `from sklearn.model_selection import train_test_split`  
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=`

In [ ]: `from sklearn.preprocessing import StandardScaler`  
`scaler = StandardScaler()`  
`X_train_scaled = scaler.fit_transform(X_train)`

# Visualize Correlation matrix (heatmap) before and after PCA

## Before PCA

```
In [ ]: import matplotlib.pyplot as plt
import seaborn as sb
```

```
X_train_scaled_df = pd.DataFrame(data=X_train_scaled, columns=X_train.columns)
X_train_scaled_df
```

```
Out[ ]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8
0	1.022555	0.997851	-0.229626	-0.207385	0.234215	-0.367791	-0.064022	-0.505889	0.03060
1	0.471283	-0.205221	-0.378220	1.027544	-1.424101	-0.078380	0.126364	0.013567	-0.33755
2	1.153387	0.036558	0.495563	-0.370033	-0.500363	0.777856	-0.268414	0.632710	0.09898
3	-0.023638	-0.273682	0.612684	1.156521	1.957021	0.359664	0.750211	0.680997	-0.06809
4	-0.255590	-2.056777	1.145573	-0.283165	-0.019856	-0.617403	-0.358912	-0.351206	1.10558
...	...	...	...	...	...	...	...	...	...
227840	-0.403774	0.598760	0.061075	0.324476	0.327051	-0.214298	-0.159001	-0.133772	0.10062
227841	1.352067	-0.396713	0.087171	-0.753679	-0.875555	1.397573	2.926905	-0.375714	1.14500
227842	-0.315815	-0.075332	0.599620	1.007243	0.344127	0.251006	-0.609354	0.864946	-0.33225
227843	-0.144489	-1.506155	1.421728	-1.664055	-2.682097	1.344205	2.040924	-0.380267	1.86595
227844	-0.387707	0.629238	-0.473540	0.255988	-0.492355	-0.745868	-0.476073	-0.404876	-0.15789


227845 rows × 30 columns

```
In [ ]: X_train_scaled_df.corr()
```

Out[ ]:

	Time	V1	V2	V3	V4	V5	V6	V7	
<b>Time</b>	1.000000	0.117203	-0.010844	-0.421238	-0.104487	0.171204	-0.063108	0.083313	-0.0387
<b>V1</b>	0.117203	1.000000	0.007295	-0.000782	-0.000227	0.001902	-0.003035	-0.011855	0.0005
<b>V2</b>	-0.010844	0.007295	1.000000	0.008112	-0.003559	0.008695	-0.003240	-0.005729	0.0014
<b>V3</b>	-0.421238	-0.000782	0.008112	1.000000	-0.001850	0.002671	-0.004517	-0.014211	0.0022
<b>V4</b>	-0.104487	-0.000227	-0.003559	-0.001850	1.000000	-0.002312	0.002631	0.005144	-0.0001
<b>V5</b>	0.171204	0.001902	0.008695	0.002671	-0.002312	1.000000	-0.008561	-0.020789	0.0031
<b>V6</b>	-0.063108	-0.003035	-0.003240	-0.004517	0.002631	-0.008561	1.000000	0.010363	-0.0011
<b>V7</b>	0.083313	-0.011855	-0.005729	-0.014211	0.005144	-0.020789	0.010363	1.000000	-0.0069
<b>V8</b>	-0.038743	0.000546	0.001479	0.002246	-0.000180	0.003117	-0.001161	-0.006947	1.0000
<b>V9</b>	-0.010310	-0.001120	-0.000090	-0.001977	0.000379	-0.001800	0.000045	-0.005432	0.0005
<b>V10</b>	0.029694	-0.001860	0.004352	-0.000745	0.000909	-0.001075	0.000081	-0.011277	-0.0023
<b>V11</b>	-0.247385	-0.000448	-0.001157	0.000983	0.002178	-0.001301	0.000622	-0.000509	0.0030
<b>V12</b>	0.125590	-0.001497	-0.000544	-0.000075	-0.001699	-0.000569	0.000067	-0.001953	-0.0035
<b>V13</b>	-0.066410	-0.001924	-0.001757	-0.001095	-0.000702	-0.000614	-0.000054	0.001543	-0.0007
<b>V14</b>	-0.098385	0.000033	-0.001041	0.000494	-0.000707	0.000905	0.000656	0.001256	-0.0025
<b>V15</b>	-0.183030	-0.001857	0.000280	-0.001326	0.001253	-0.001777	0.001112	0.002246	-0.0001
<b>V16</b>	0.012980	-0.002791	0.000076	-0.001699	0.000616	-0.002415	0.002639	0.005372	-0.0041
<b>V17</b>	-0.071968	-0.001581	0.000843	-0.000341	0.000288	0.000645	-0.000140	-0.002626	-0.0046
<b>V18</b>	0.091128	-0.000757	0.000502	-0.001566	0.001995	0.000903	0.000657	0.000283	-0.0026
<b>V19</b>	0.029596	0.001869	0.001993	0.001622	-0.001183	-0.001521	0.000371	-0.002998	-0.0006
<b>V20</b>	-0.051246	-0.001306	-0.004863	0.001906	0.000714	0.005301	-0.002250	-0.000800	0.0015
<b>V21</b>	0.045097	-0.001346	0.002492	-0.000554	0.001492	0.001764	-0.001328	-0.008404	-0.0064
<b>V22</b>	0.144891	0.003286	0.001870	0.000905	0.000647	-0.000594	-0.000003	0.000382	0.0037
<b>V23</b>	0.050655	0.005771	0.004074	0.004907	-0.002682	0.001927	-0.001865	0.000138	0.0025
<b>V24</b>	-0.015464	-0.001356	-0.000414	-0.000579	-0.000989	-0.001035	0.000782	0.000365	-0.0002
<b>V25</b>	-0.232573	0.001797	0.001938	0.000258	-0.001542	-0.000433	0.000644	0.000310	-0.0018
<b>V26</b>	-0.042594	-0.000779	0.000773	-0.000720	0.001766	-0.000076	-0.000277	0.001284	-0.0011
<b>V27</b>	-0.005570	-0.002175	-0.000278	-0.003023	0.001041	-0.008595	0.004883	0.008366	-0.0058
<b>V28</b>	-0.009371	-0.003073	-0.004816	0.002885	0.002495	0.008496	-0.005313	-0.006193	0.0011
<b>Amount</b>	-0.009936	-0.233925	-0.536033	-0.218054	0.103808	-0.394113	0.221781	0.411463	-0.1054


30 rows × 30 columns



```
In [ ]: plt.figure(figsize=(25, 25))

sb.heatmap(X_train_scaled_df.corr(), annot=True)
```

Out[ ]: <Axes: >



## After PCA

```
In [ ]: from sklearn.decomposition import PCA
pca = PCA()
X_train_PCA = pca.fit_transform(X_train_scaled)
X_train_PCA
```

```
Out[ ]: array([[ -0.45171765, -1.29559725,  0.0421551 , ...,  0.03195231,
          -0.13785439,  0.01672698],
        [ -0.25134876, -1.31061702, -0.10325889, ..., -0.27767841,
          0.65829119,  0.02265744],
        [ -0.40706699, -1.42453283,  0.32045815, ..., -0.28999978,
          -0.18890148,  0.04057583],
        ...,
        [ -0.4183145 ,  0.61626946,  0.26434248, ..., -0.6692277 ,
          -0.1581937 ,  0.02980141],
        [ -0.58662483, -0.21794326,  0.24625923, ...,  0.69319465,
          0.4484557 , -0.11181253],
        [  0.12975006,  0.74715442,  0.07854037, ...,  0.04452226,
          -0.20418901, -0.01583862]])
```

```
In [ ]: X_train_PCA_df = pd.DataFrame(data=X_train_PCA, columns = X_train.columns)
X_train_PCA_df
```

```
Out[ ]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8
0	-0.451718	-1.295597	0.042155	-0.017005	-0.558668	-0.103815	-0.540228	-0.405253	-1.88707
1	-0.251349	-1.310617	-0.103259	0.402708	-0.246232	0.388712	-1.104457	1.130561	-0.90238
2	-0.407067	-1.424533	0.320458	0.460765	0.406264	0.957740	-0.135362	0.880158	0.39857
3	0.195144	0.444337	0.441017	-0.301536	-0.288001	2.739601	-1.813253	-0.484622	0.27349
4	-0.496166	0.505630	0.756510	0.071279	0.729322	-0.565272	0.492185	-0.453229	-1.04485
...	...	...	...	...	...	...	...	...	...
227840	-0.432766	1.085389	0.103383	0.134239	0.101425	-0.207721	-0.077260	-0.501031	-0.43097
227841	-0.050833	-2.107870	0.347575	1.184980	-0.696713	-0.527504	0.206993	0.797536	-0.69992
227842	-0.418314	0.616269	0.264342	-0.615669	-0.494755	-0.551815	0.203305	1.002262	0.20252
227843	-0.586625	-0.217943	0.246259	0.750738	0.238745	-1.475042	-0.287662	0.415580	-0.01622
227844	0.129750	0.747154	0.078540	-0.175870	0.115263	-0.629163	0.838574	0.687421	0.54297

227845 rows × 30 columns

```
In [ ]: X_train_PCA_df.corr()
```

Out[ ]:

	Time	V1	V2	V3	V4	V5	
<b>Time</b>	1.000000e+00	7.839378e-17	1.708444e-16	-1.868911e-16	-7.130918e-17	-1.278624e-18	1.850
<b>V1</b>	7.839378e-17	1.000000e+00	1.586145e-17	4.744220e-17	3.927787e-17	1.291499e-17	1.004
<b>V2</b>	1.708444e-16	1.586145e-17	1.000000e+00	-3.124562e-16	2.323815e-16	6.032887e-17	1.374
<b>V3</b>	-1.868911e-16	4.744220e-17	-3.124562e-16	1.000000e+00	3.768324e-16	2.144041e-17	-6.7
<b>V4</b>	-7.130918e-17	3.927787e-17	2.323815e-16	3.768324e-16	1.000000e+00	9.352089e-17	-4.6
<b>V5</b>	-1.278624e-18	1.291499e-17	6.032887e-17	2.144041e-17	9.352089e-17	1.000000e+00	-2.5
<b>V6</b>	1.850432e-17	1.004389e-17	1.374008e-17	-6.739788e-17	-4.612447e-16	-2.591059e-16	1.0000
<b>V7</b>	-3.954003e-18	-3.221124e-17	1.924877e-17	1.474727e-17	1.499357e-17	5.091816e-17	-2.4
<b>V8</b>	9.460095e-17	4.235341e-18	-3.849014e-17	-3.222983e-17	3.990450e-17	6.472387e-16	-7.2
<b>V9</b>	-5.799076e-19	-1.543606e-17	-8.561077e-18	-8.950723e-17	-4.318201e-17	1.124380e-16	1.890
<b>V10</b>	1.376964e-16	-2.502082e-17	5.300461e-17	-4.520668e-16	-9.448290e-17	-3.576772e-18	-4.1
<b>V11</b>	-3.046093e-19	2.671807e-17	-4.289011e-16	8.111923e-17	4.650798e-17	1.806107e-16	4.113
<b>V12</b>	-1.010646e-16	-1.730079e-17	4.223591e-17	-3.273198e-17	-1.405945e-16	2.476665e-16	-1.6
<b>V13</b>	5.161963e-17	9.557533e-18	-1.139906e-16	3.559297e-17	-5.964028e-17	1.548470e-16	6.481
<b>V14</b>	1.097591e-16	1.582437e-17	7.103984e-18	-1.620744e-16	-2.064736e-17	2.188318e-17	-9.6
<b>V15</b>	-2.124862e-17	2.554684e-17	2.258294e-17	2.283802e-17	-2.277063e-17	1.145358e-16	4.835
<b>V16</b>	-7.528998e-17	5.555474e-17	2.558791e-17	-1.197628e-16	-1.729398e-16	-2.073014e-16	-1.7
<b>V17</b>	1.561661e-16	6.777082e-17	2.066603e-16	1.450569e-17	1.242875e-16	-1.331043e-16	-8.0
<b>V18</b>	-2.146808e-17	1.104800e-17	-8.919885e-17	-2.533323e-16	-1.955507e-16	1.068794e-16	-5.1
<b>V19</b>	-7.599261e-17	1.750763e-17	2.561093e-18	-2.465785e-16	-2.236296e-16	1.954372e-16	-7.5
<b>V20</b>	5.734832e-17	-4.964350e-18	1.024628e-16	1.043208e-16	1.777355e-16	8.674517e-18	-1.4


	Time	V1	V2	V3	V4	V5	
<b>V21</b>	-2.409424e-17	5.719831e-18	-6.970960e-17	1.091305e-16	2.194394e-16	-3.663227e-16	-6.9
<b>V22</b>	-6.295647e-17	-2.254200e-17	1.148360e-16	-8.685002e-17	1.316876e-16	3.648188e-16	2.540
<b>V23</b>	1.446412e-16	8.123133e-17	-1.049816e-16	-7.088991e-17	4.172204e-17	-1.880598e-16	6.161
<b>V24</b>	7.828680e-18	3.572490e-19	8.026503e-17	4.922013e-17	2.317600e-16	1.434056e-16	4.123
<b>V25</b>	-1.392130e-17	2.766950e-17	-5.555925e-18	2.393508e-16	4.345410e-17	-1.597242e-16	1.873
<b>V26</b>	1.799874e-16	2.429678e-17	4.658041e-16	-2.328068e-16	-1.715146e-16	1.783046e-17	-1.7
<b>V27</b>	6.937767e-17	2.880595e-18	-3.584702e-16	-7.615352e-16	1.153877e-16	5.038011e-17	1.191
<b>V28</b>	-2.641047e-17	-4.021459e-16	3.332590e-17	2.806540e-17	-4.258394e-17	1.372546e-17	1.616
<b>Amount</b>	3.141768e-16	-5.310173e-18	-4.573220e-16	1.104138e-16	-3.666818e-16	-7.347193e-17	5.878

30 rows × 30 columns

```
In [ ]: plt.figure(figsize=(25, 25))
        sb.heatmap(X_train_PCA_df.corr(), annot=True)
```

Out[ ]: <Axes: >



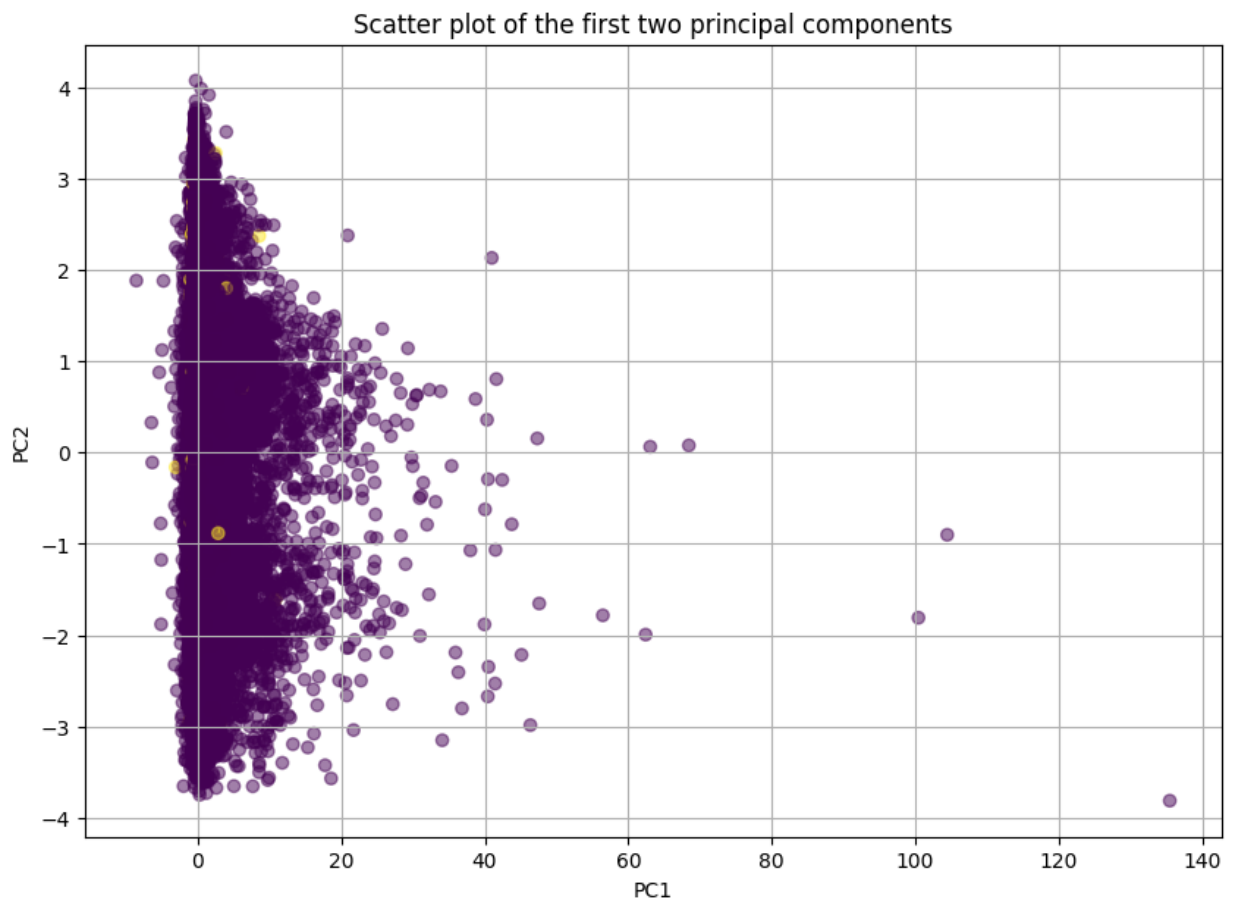


## Visualize the classes in 2D. Discuss corner cases.

```
In [ ]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 7))

plt.scatter(X_train_PCA_df.iloc[:, 0], X_train_PCA_df.iloc[:, 1], c=y_train, cmap='viridis')
plt.xlabel('PC1')
plt.ylabel('PC2')
plt.title('Scatter plot of the first two principal components')
plt.grid(True)
plt.show()
```



## Apply SVC on Model 1, raw

```
In [ ]: from sklearn.svm import SVC
        from sklearn.metrics import roc_curve, auc
        from sklearn.preprocessing import label_binarize

        model = SVC(probability=True)
        model.fit(X_train_scaled, y_train)
```

```
Out[ ]: SVC
        SVC(probability=True)
```

```
In [ ]: X_test_scaled = scaler.transform(X_test)

        y_score = model.predict_proba(X_test_scaled)
        y_score
```

```
Out[ ]: array([[2.92213182e-01, 7.07786818e-01],
               [9.99546134e-01, 4.53865621e-04],
               [9.99447428e-01, 5.52572344e-04],
               ...,
               [9.99653336e-01, 3.46663971e-04],
               [9.99671487e-01, 3.28513334e-04],
               [9.99248522e-01, 7.51477857e-04]])
```

```
In [ ]: y_score_m1 = model.predict_proba(X_test_scaled)[: ,1]
        fpr_m1, tpr_m1, _ = roc_curve(y_test, y_score_m1)
        roc_auc_m1 = auc(fpr_m1, tpr_m1)
```

## M2 transformed data. (all features (PCs))

```
In [ ]: model_2 = SVC(probability=True)
        model_2.fit(X_train_PCA, y_train)
```

```
Out[ ]: SVC
        SVC(probability=True)
```

```
In [ ]: X_test_pca = pca.transform(X_test_scaled)

        y_score_m2 = model_2.predict_proba(X_test_pca)[: , 1]
        fpr_m2, tpr_m2, _ = roc_curve(y_test, y_score_m2)
        roc_auc_m2 = auc(fpr_m2, tpr_m2)
```

## M3-6 reduced data from (top, 1, 2, 3, 4, 5 features (PCs))

```
In [ ]: models_m3_to_m7 = []

        for i in range(1, 6):
            svc = SVC(probability=True)
            svc.fit(X_train_PCA[:, :i], y_train)

            X_test_PCA_reduced = pca.transform(scaler.transform(X_test))[:, :i]

            y_score = svc.predict_proba(X_test_PCA_reduced)[: , 1]

            fpr, tpr, _ = roc_curve(y_test, y_score)
            roc_auc = auc(fpr, tpr)

            models_m3_to_m7.append({
                'fpr': fpr,
                'tpr': tpr,
                'roc_auc': roc_auc,
                'label': f'M{i+2} (Top {i} PCs)'
            })
```

## Report the findings for all six models in terms of ROC curves on one plot.

```
In [ ]: plt.figure(figsize=(10, 8))

        plt.plot(fpr_m1, tpr_m1, label=f'M1 (Raw Data) (area = {roc_auc_m1:.2f})')

        plt.plot(fpr_m2, tpr_m2, label=f'M2 (All Features - PCA) (area = {roc_auc_m2:.2f})')
```


```

for model in models_m3_to_m7:
    plt.plot(model['fpr'], model['tpr'], label=f"{model['label']} (area = {model['roc_

plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curves for Different Models')
plt.legend(loc="best")

plt.show()

```



## Discuss the merits/de-merits of each model.

### Model 1:

**Strengths:** The model works really well with the untouched data, showing that the original details are quite helpful and a good fit for sorting or categorizing.

**Weaknesses:** Keeping all the original details could make the model more complicated and might lead to it being too tailored to the training data (overfitting). This could also mean it takes more time to learn from the data and to make predictions, especially with a big dataset.

#### Model 2:

Strengths: It works as well as Model 1, showing that the PCA method kept most of the important changes in the data. It could also be better at handling random variations and differences in the data than Model 1.

Weaknesses: Like Model 1, using all the components doesn't simplify the model. It might be harder to understand how the model makes decisions because the PCA components don't directly match up with the original details.

#### Model 3:

Strengths: The model is straightforward and quick to train because it uses only the most important principal component.

Weaknesses: There's a big decrease in how well it performs compared to Models 1 and 2, indicating that relying solely on the top principal component isn't sufficient for precise categorization.

#### Model 4:

Strengths: It's a bit more complex than Model 3 but remains simpler and quicker than Models 1 and 2.

Weaknesses: The model's AUC (Area Under the Curve) is below 0.5, which means it performs worse than if it were just making random guesses. This suggests that including the top two principal components might add confusing patterns or noise to the model, leading to poor predictions.

#### Model 5:

Strengths: Using three principal components is more informative than using just one or two, and it also simplifies the model more than using all the original features.

Weaknesses: The model's performance is much worse than random guessing, as indicated by an AUC (Area Under the Curve) significantly lower than 0.5. This suggests that the top three principal components may not be capturing the necessary information for this specific classification task.

#### Model 6:

Strengths: There's a small improvement compared to Models 4 and 5, which implies that adding the fourth principal component provides some useful information for classifying.

Weaknesses: The model's effectiveness is still close to what you'd expect from random guesses, showing that the top four principal components don't provide enough detail for strong classification results.

#### Model 7:

Strengths: There's a notable boost in how well it works, nearly matching the levels of Models 1 and 2. This indicates that using the top five principal components gets most of the crucial details needed for sorting or categorizing, and adding just a few more details could further help.

Weaknesses: Although it outperforms Models 3 to 6, it's still not quite as precise as Models 1 and 2. This suggests there could be a balance to find between keeping the model simple and achieving the best possible performance.