

```
In [ ]: import torch
from torchvision import datasets, transforms
from torch.utils.data import DataLoader
from sklearn.decomposition import PCA
from sklearn.metrics import mean_squared_error
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from matplotlib import pyplot as plt
```

## Importing Data

```
In [ ]: # Transform to convert images to PyTorch tensors and normalize them
transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,)) # Normalize around the mean 0.5
])

# Download and Load the training data
trainset = datasets.MNIST(root='./data', download=True, train=True, transform=transform)
trainloader = DataLoader(trainset, batch_size=64, shuffle=True)

# Download and Load the test data
testset = datasets.MNIST(root='./data', download=True, train=False, transform=transform)
testloader = DataLoader(testset, batch_size=64, shuffle=True)
```

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz  
 Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to ./data/MNIST/raw/train-images-idx3-ubyte.gz

100%|██████████| 9912422/9912422 [00:00<00:00, 219483760.03it/s]

Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz  
 Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to ./data/MNIST/raw/train-labels-idx1-ubyte.gz

100%|██████████| 28881/28881 [00:00<00:00, 49503757.18it/s]

Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz  
 Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw/t10k-images-idx3-ubyte.gz

100%|██████████| 1648877/1648877 [00:00<00:00, 78595926.91it/s]

Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz  
 Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz

100%|██████████| 4542/4542 [00:00<00:00, 16638016.39it/s]

Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw

## PCA

```
In [ ]: # Flatten the images for PCA
X_train_flattened = trainset.data.view(len(trainset), -1).numpy()
X_test_flattened = testset.data.view(len(testset), -1).numpy()

# Normalize data
X_train_flattened = X_train_flattened / 255.0
X_test_flattened = X_test_flattened / 255.0

# PCA with different number of components
components = [50, 100, 150, 200]
for n_components in components:
    pca = PCA(n_components=n_components)
    pca.fit(X_train_flattened)
    X_train_pca = pca.transform(X_train_flattened)
    X_test_pca = pca.transform(X_test_flattened)
    X_reconstructed_pca = pca.inverse_transform(X_train_pca)
    reconstruction_loss = mean_squared_error(X_train_flattened, X_reconstructed_pca)

# Logistic regression classifier to classify our reconstructed images
lr = LogisticRegression(max_iter=1000, solver='sag', multi_class='multinomial')
lr.fit(X_train_pca, trainset.targets.numpy())

# Evaluating on a subset of the test set
validation_index = np.random.choice(len(testset), 1000, replace=False)
validation_images_pca = X_test_pca[validation_index]
validation_labels = testset.targets.numpy()[validation_index]

validation_predictions = lr.predict(validation_images_pca)
accuracy = accuracy_score(validation_labels, validation_predictions)
print(f'Reconstruction loss for PCA with {n_components} components: {reconstruction_loss} Accuracy: {accuracy}')
```

Reconstruction loss for PCA with 50 components: 0.011795084610364161 Classification Accuracy: 0.924

Reconstruction loss for PCA with 100 components: 0.005756633405031867 Classification Accuracy: 0.93

Reconstruction loss for PCA with 150 components: 0.0034891502018314425 Classification Accuracy: 0.924


Reconstruction loss for PCA with 200 components: 0.0022767954728689765 Classification Accuracy: 0.926

The best scoring minimal components is 100 as it has the highest accuracy and a decent reconstruction loss.

```
In [ ]: # Visualizing the first 25 PCA'S
fig, axes = plt.subplots(5, 5, figsize=(12, 10))

for i in range(25):
    row = i // 5
    col = i % 5
    pc_image = pca.components_[i].reshape(28, 28) # Reshape the principal component to 28x28
    axes[row, col].imshow(pc_image, cmap='gray')
    axes[row, col].set_title(f'PC {i+1}')
    axes[row, col].axis('off')

plt.tight_layout()
plt.show()
```



```
In [ ]: # Comparing the reconstructed images to the original images
fig, axes = plt.subplots(5, 2, figsize=(8, 4 * 5))

for i in range(5):
    # Original image
    original_image = X_train_flattened[i].reshape(28, 28)


    # Reconstructed image after PCA
    reconstructed_image = pca.inverse_transform(X_train_pca[i]).reshape(28, 28)

    # Plot original image
    axes[i, 0].imshow(original_image, cmap='gray')
    axes[i, 0].set_title(f'Original Image {i+1}')
    axes[i, 0].axis('off')


    # Plot reconstructed image
    axes[i, 1].imshow(reconstructed_image, cmap='gray')
    axes[i, 1].set_title(f'Reconstructed Image {i+1}')
    axes[i, 1].axis('off')

plt.tight_layout()
plt.show()
```


Original Image 1




Reconstructed Image 1




Original Image 2




Reconstructed Image 2



Original Image 3



Reconstructed Image 3




Original Image 4




Reconstructed Image 4




Original Image 4




Reconstructed Image 4



Original Image 5



Reconstructed Image 5




```
In [ ]: def generate_digits(n_digits):
fig, axes = plt.subplots(1, n_digits, figsize=(10, 2))
for i in range(n_digits):
    # Randomly sample from the distribution of the PCA coefficients
    random_sample = np.random.normal(0, 1, pca.n_components)

    # Generate new data by inverse transforming the random sample
    new_digit = pca.inverse_transform(random_sample)

    # Rescale to original range and reshape to 28x28
    new_digit_image = new_digit.reshape(28, 28)

    # Plot
    ax = axes[i]
    ax.imshow(new_digit_image, cmap='gray')
    ax.axis('off')
plt.show()
```

```
# Generate and display 5 new digits
generate_digits(5)
```



The images were generated by selecting random points within a space transformed by Principal Component Analysis (PCA) and then converting these points back into the format of the original images. As a result, these images might not look exactly like any specific digits from the MNIST dataset. Instead, they display variations based on the major trends or features identified by the PCA.

## Autoencoder

```
In [ ]: import torch.nn as nn
import torch.optim as optim

class Autoencoder(nn.Module):
    def __init__(self, input_size, hidden_size):
        super(Autoencoder, self).__init__()
        self.encoder = nn.Linear(input_size, hidden_size)
        self.decoder = nn.Linear(hidden_size, input_size)

    def forward(self, x):
        x = torch.flatten(x, start_dim=1)
        encoded = torch.relu(self.encoder(x))
        decoded = torch.sigmoid(self.decoder(encoded))
        return decoded

# Train autoencoder
def train_autoencoder(hidden_size, results={}):

    model = Autoencoder(784, hidden_size)
    criterion = nn.MSELoss()
    optimizer = optim.Adam(model.parameters(), lr=1e-3)

    # Training Loop
    num_epochs = 8
    for epoch in range(num_epochs):
        total_loss = 0
        for images, _ in trainloader:
            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, torch.flatten(images, start_dim=1))
            loss.backward()
            optimizer.step()
            total_loss += loss.item()
        average_loss = total_loss / len(trainloader)

    results[hidden_size] = {'reconstruction_loss': average_loss}
```

```

encoded_train = []
train_labels = []
for data, target in trainloader:
    encoded = model(data)
    encoded_train.append(encoded.detach().cpu().numpy())
    train_labels.append(target.numpy())
encoded_train = np.concatenate(encoded_train, axis=0)
train_labels = np.concatenate(train_labels, axis=0)

clf = LogisticRegression(max_iter=1000, solver='sag', multi_class='multinomial')
clf.fit(encoded_train, train_labels)

encoded_val = []
val_labels = []
for data, target in testloader:
    encoded = model(data)
    encoded_val.append(encoded.detach().cpu().numpy())
    val_labels.append(target.numpy())
encoded_val = np.concatenate(encoded_val, axis=0)
val_labels = np.concatenate(val_labels, axis=0)
val_predictions = clf.predict(encoded_val)
accuracy = accuracy_score(val_labels, val_predictions)

results[hidden_size]['classification_accuracy'] = accuracy
print(f'Encoding Dim: {hidden_size}, Classification Accuracy: {accuracy}')

return model, loss.item(), results

hidden_sizes = [50, 100, 150, 200]
results = {}
for hidden_size in hidden_sizes:
    model, loss, results = train_autoencoder(hidden_size, results)
    print(f'Reconstruction loss for Autoencoder with {hidden_size} hidden units:', loss)
for encoding_dim, metrics in results.items():
    print(f'{encoding_dim}: Reconstruction Loss: {metrics["reconstruction_loss"]}, Cla

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: Convergence
Warning: The max_iter was reached which means the coef_ did not converge

```

```

warnings.warn(
Encoding Dim: 50, Classification Accuracy: 0.8038
Reconstruction loss for Autoencoder with 50 hidden units: 0.9075490236282349

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: Convergence
Warning: The max_iter was reached which means the coef_ did not converge

```

```

warnings.warn(
Encoding Dim: 100, Classification Accuracy: 0.8621
Reconstruction loss for Autoencoder with 100 hidden units: 0.8906221985816956

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: Convergence
Warning: The max_iter was reached which means the coef_ did not converge

```

```

warnings.warn(
Encoding Dim: 150, Classification Accuracy: 0.8531
Reconstruction loss for Autoencoder with 150 hidden units: 0.8868787884712219

```

```

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_sag.py:350: Convergence
Warning: The max_iter was reached which means the coef_ did not converge

```

```

warnings.warn(

```

Encoding Dim: 200, Classification Accuracy: 0.8583  
 Reconstruction loss for Autoencoder with 200 hidden units: 0.883091390132904  
 50: Reconstruction Loss: 0.9358660034152236, Classification Accuracy: 0.8038  
 100: Reconstruction Loss: 0.9164198400623509, Classification Accuracy: 0.8621  
 150: Reconstruction Loss: 0.9233703469035468, Classification Accuracy: 0.8531  
 200: Reconstruction Loss: 0.9157706576623896, Classification Accuracy: 0.8583


The best minimum hidden encoding dimension was 100 because it had the highest classification accuracy while still having the 2nd lowest reconstruction loss.

```
In [ ]: hidden_size = 100
autoencoder = Autoencoder(784, hidden_size)
criterion = nn.MSELoss()
optimizer = optim.Adam(autoencoder.parameters(), lr=0.001)

for epoch in range(8):
    for data, _ in trainloader:
        optimizer.zero_grad()
        output = autoencoder(data)
        loss = criterion(output, torch.flatten(data, start_dim=1))
        loss.backward()
        optimizer.step()

# Generate new handwritten digits using the trained autoencoder
np.random.seed(42)
random_noise = np.random.normal(0,1,(30, hidden_size))
random_noise = torch.tensor(random_noise, dtype=torch.float32)
generated_images = autoencoder.decoder(random_noise).detach().cpu().numpy()

# Plot the generated images
fig, axes = plt.subplots(3, 10, figsize=(15, 4.5)) # Adjusted for 3 rows and 10 columns
for i, ax in enumerate(axes.flatten()):
    ax.imshow(generated_images[i].reshape(28, 28), cmap='gray')
    ax.axis('off')
plt.tight_layout()
plt.show()
```



## RBN

```
In [11]: # Load the MNIST dataset
transform = transforms.Compose([transforms.ToTensor(),
                                transforms.Normalize((0.5,), (0.5,))])
train = datasets.MNIST(root='./data', train=True, download=True, transform=transform)
test = datasets.MNIST(root='./data', train=False, download=True, transform=transform)
```



```
trainloader = torch.utils.data.DataLoader(train, batch_size=64, shuffle=True)
testloader = torch.utils.data.DataLoader(test, batch_size=64, shuffle=False)
```

```
In [ ]: # Define the RBM class
class RBM(torch.nn.Module):
    def __init__(self, input, hiddensize):
        super(RBM, self).__init__()
        self.W = torch.nn.Parameter(torch.randn(hiddensize, input) * 0.01)
        self.v_bias = torch.nn.Parameter(torch.zeros(input))
        self.h_bias = torch.nn.Parameter(torch.zeros(hiddensize))

    def sample_h(self, v):
        h_prob = torch.sigmoid(torch.matmul(v, self.W.t()) + self.h_bias)
        h_sample = torch.bernoulli(h_prob)
        return h_sample

    def sample_v(self, h):
        v_prob = torch.sigmoid(torch.matmul(h, self.W) + self.v_bias)
        v_sample = torch.bernoulli(v_prob)
        return v_sample

    def free_energy(self, v):
        vbias_term = torch.matmul(v, self.v_bias)
        wx_b = torch.matmul(v, self.W.t()) + self.h_bias
        hidden_term = torch.sum(torch.log(1 + torch.exp(wx_b)), dim=1)
        return -vbias_term - hidden_term

    def forward(self, v):
        h = self.sample_h(v)
        v_recon = self.sample_v(h)
        return v_recon

# List of encoding dimensions to test
hidden_size = [50, 100, 150, 200]
results_dict = {}

for size in hidden_size:
    rbm = RBM(input=784, hiddensize=size)
    optimizer = torch.optim.SGD(rbm.parameters(), lr=0.1)

    # Training Loop for RBM
    for epoch in range(8): # Adjust the number of epochs if necessary
        loss_sum = 0
        for data, _ in trainloader:
            data = data.view(-1, 784) # Flatten the images
            data = torch.clamp(data, 0.0, 1.0)
            data = (data > 0.5).float() # Convert to binary using thresholding
            v_recon = rbm(data)
            loss = torch.mean(rbm.free_energy(data)) - torch.mean(rbm.free_energy(v_recon))
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()
            loss_sum += loss.item()
        average_loss = loss_sum / len(trainloader)
        results_dict[size] = {'reconstruction_loss': average_loss}

# Extract features from the training set
encoded_train = []
train_labels = []
```

```

for data, target in trainloader:
    data = data.view(-1, 784)
    data = torch.clamp(data, 0.0, 1.0)
    data = (data > 0.5).float() # Convert to binary using thresholding
    encoded = rbm.sample_h(data) # Get encoded (hidden) representations
    encoded_train.append(encoded.detach().numpy())
    train_labels.append(target.numpy())
encoded_train = np.concatenate(encoded_train, axis=0)
train_labels = np.concatenate(train_labels, axis=0)

# Train a softmax classifier
lr = LogisticRegression(max_iter=1000, solver='lbfgs', multi_class='multinomial')
lr.fit(encoded_train, train_labels)

# Evaluate classification accuracy on the test set
encoded_test = []
test_labels = []
for data, target in testloader:
    data = data.view(-1, 784)
    data = torch.clamp(data, 0.0, 1.0)
    data = (data > 0.5).float() # Convert to binary using thresholding
    encoded = rbm.sample_h(data)
    encoded_test.append(encoded.detach().numpy())
    test_labels.append(target.numpy())
encoded_test = np.concatenate(encoded_test, axis=0)
test_labels = np.concatenate(test_labels, axis=0)
test_predictions = lr.predict(encoded_test)
accuracy = accuracy_score(test_labels, test_predictions)

# Record the classification accuracy
results_dict[size]['classification_accuracy'] = accuracy
print(f'Hidden Size: {size}, Reconstruction Loss: {average_loss}, Classification A

# Print the results for all encoding dimensions
print("Results for different Hidden Size:")
for size, metrics in results_dict.items():
    print(f'{size}: Reconstruction Loss: {metrics["reconstruction_loss"]}, Classificat

```

```

Encoding Dim: 50, Reconstruction Loss: -19.923210117608498, Classification Accuracy:
0.8439
Encoding Dim: 75, Reconstruction Loss: -18.55157110787658, Classification Accuracy:
0.8666
Encoding Dim: 100, Reconstruction Loss: -17.291857076860442, Classification Accuracy:
0.8758
Encoding Dim: 150, Reconstruction Loss: -15.531131760652132, Classification Accuracy:
0.8845
Encoding Dim: 200, Reconstruction Loss: -14.382853977715792, Classification Accuracy:
0.8889
Results for different Encoding Dimensions:
50: Reconstruction Loss: -19.923210117608498, Classification Accuracy: 0.8439
75: Reconstruction Loss: -18.55157110787658, Classification Accuracy: 0.8666
100: Reconstruction Loss: -17.291857076860442, Classification Accuracy: 0.8758
150: Reconstruction Loss: -15.531131760652132, Classification Accuracy: 0.8845
200: Reconstruction Loss: -14.382853977715792, Classification Accuracy: 0.8889

```

The best minimum hidden encoding dimension was 150 because it had the 2nd highest classification and the 2nd lowest reconstruction loss but it is about the same as our max value which is 200. When trying to go for efficiency and the smallest size we can go, 150 is better than 200.


```
In [ ]: # Set the number of hidden dimensions
hiddensize = 150

# Train an RBM with 150 hidden dimensions (assuming train_loader is defined)
rbm = RBM(n_vis=784, n_hid=hiddensize)
optimizer = torch.optim.SGD(rbm.parameters(), lr=0.1)

for epoch in range(8): # Adjust the number of epochs if necessary
    for data, _ in trainloader:
        data = data.view(-1, 784)
        data = torch.clamp(data, 0.0, 1.0)
        data = (data > 0.5).float() # Convert to binary using thresholding
        v_recon = rbm(data)
        loss = torch.mean(rbm.free_energy(data)) - torch.mean(rbm.free_energy(v_recon))
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

# Generate 30 new images using the trained RBM
np.random.seed(42) # Set a seed for reproducibility
random_hidden = torch.bernoulli(torch.rand(30, hiddensize)) # Randomly sample the hidden states
generated_images = rbm.sample_v(random_hidden).detach().numpy() # Reconstruct the visible states

# Plot the generated images
fig, axes = plt.subplots(3, 10, figsize=(15, 4.5)) # Adjusted for 3 rows and 10 columns
for i, ax in enumerate(axes.flatten()):
    ax.imshow(generated_images[i].reshape(28, 28), cmap='gray')
    ax.axis('off')
plt.tight_layout()
plt.show()
```



## DBN

```
In [4]: !pip install keras-tuner
```

Collecting keras-tuner

Downloading keras\_tuner-1.4.7-py3-none-any.whl (129 kB)

129.1/129.1 kB 1.9 MB/s eta 0:00:00

Requirement already satisfied: keras in /usr/local/lib/python3.10/dist-packages (from keras-tuner) (3.2.1)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from keras-tuner) (24.0)

Requirement already satisfied: requests in /usr/local/lib/python3.10/dist-packages (from keras-tuner) (2.31.0)

Collecting kt-legacy (from keras-tuner)

Downloading kt\_legacy-1.0.5-py3-none-any.whl (9.6 kB)

Requirement already satisfied: absl-py in /usr/local/lib/python3.10/dist-packages (from keras->keras-tuner) (1.4.0)

Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from keras->keras-tuner) (1.25.2)

Requirement already satisfied: rich in /usr/local/lib/python3.10/dist-packages (from keras->keras-tuner) (13.7.1)

Requirement already satisfied: namex in /usr/local/lib/python3.10/dist-packages (from keras->keras-tuner) (0.0.7)

Requirement already satisfied: h5py in /usr/local/lib/python3.10/dist-packages (from keras->keras-tuner) (3.9.0)

Requirement already satisfied: optree in /usr/local/lib/python3.10/dist-packages (from keras->keras-tuner) (0.11.0)

Requirement already satisfied: ml-dtypes in /usr/local/lib/python3.10/dist-packages (from keras->keras-tuner) (0.2.0)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests->keras-tuner) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests->keras-tuner) (3.6)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests->keras-tuner) (2.0.7)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests->keras-tuner) (2024.2.2)

Requirement already satisfied: typing-extensions>=4.0.0 in /usr/local/lib/python3.10/dist-packages (from optree->keras->keras-tuner) (4.11.0)

Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.10/dist-packages (from rich->keras->keras-tuner) (3.0.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.10/dist-packages (from rich->keras->keras-tuner) (2.16.1)

Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-it-py>=2.2.0->rich->keras->keras-tuner) (0.1.2)

Installing collected packages: kt-legacy, keras-tuner

Successfully installed keras-tuner-1.4.7 kt-legacy-1.0.5

```
In [7]: import tensorflow as tf
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.models import Sequential
from tensorflow.keras.datasets import mnist
from tensorflow.keras.utils import to_categorical
from kerastuner.tuners import RandomSearch

def load_data():
    # Load and preprocess the MNIST dataset
    (x_train, y_train), (x_test, y_test) = mnist.load_data()
    x_train = x_train.reshape(-1, 784).astype('float32') / 255
    x_test = x_test.reshape(-1, 784).astype('float32') / 255
    y_train = to_categorical(y_train, 10)
    y_test = to_categorical(y_test, 10)
    return x_train, y_train, x_test, y_test
```

```

def build_model(hp):
    model = Sequential()
    # Hyperparameters to tune
    units = hp.Choice('units', values=[50, 100, 150, 200])
    dropout_rate = hp.Float('dropout_rate', min_value=0.0, max_value=0.5, step=0.1)

    # Model architecture
    model.add(Dense(units=units, activation='relu', input_shape=(784,)))
    model.add(Dropout(dropout_rate))
    model.add(Dense(units=units, activation='relu'))
    model.add(Dropout(dropout_rate))
    model.add(Dense(10, activation='softmax'))

    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    return model

def main():
    x_train, y_train, x_test, y_test = load_data()
    tuner = RandomSearch(
        build_model,
        objective='val_accuracy',
        max_trials=10, # Adjust the number of trials to cover the grid
        executions_per_trial=1,
        directory='tuner_dir',
        project_name='DBN_grid_search')

    tuner.search(x_train, y_train, epochs=10, validation_data=(x_test, y_test))

    # Get the optimal hyperparameters
    best_hps = tuner.get_best_hyperparameters()[0]
    print(f"Best number of units: {best_hps.get('units')}")
    print(f"Best dropout rate: {best_hps.get('dropout_rate')}")

    # Build the model with the best hyperparameters and train it
    model = tuner.hypermodel.build(best_hps)
    model.fit(x_train, y_train, epochs=50, validation_data=(x_test, y_test))

if __name__ == "__main__":
    main()

```

Trial 10 Complete [00h 02m 04s]  
val\_accuracy: 0.9815000295639038


Best val\_accuracy So Far: 0.9815000295639038

Total elapsed time: 00h 17m 30s


Best number of units: 200

Best dropout rate: 0.1


Epoch 1/50

1875/1875  10s 5ms/step - accuracy: 0.8755 - loss: 0.4115 - val\_accuracy: 0.9603 - val\_loss: 0.1234


Epoch 2/50

1875/1875  10s 5ms/step - accuracy: 0.9654 - loss: 0.1124 - val\_accuracy: 0.9710 - val\_loss: 0.0889


Epoch 3/50

1875/1875  11s 6ms/step - accuracy: 0.9755 - loss: 0.0782 - val\_accuracy: 0.9777 - val\_loss: 0.0756


Epoch 4/50

1875/1875  10s 5ms/step - accuracy: 0.9810 - loss: 0.0575 - val\_accuracy: 0.9805 - val\_loss: 0.0625


Epoch 5/50

1875/1875  9s 4ms/step - accuracy: 0.9847 - loss: 0.0481 - val\_accuracy: 0.9744 - val\_loss: 0.0827


Epoch 6/50

1875/1875  10s 5ms/step - accuracy: 0.9865 - loss: 0.0425 - val\_accuracy: 0.9797 - val\_loss: 0.0668


Epoch 7/50

1875/1875  10s 5ms/step - accuracy: 0.9885 - loss: 0.0359 - val\_accuracy: 0.9804 - val\_loss: 0.0721


Epoch 8/50

1875/1875  10s 5ms/step - accuracy: 0.9894 - loss: 0.0323 - val\_accuracy: 0.9797 - val\_loss: 0.0726

Epoch 9/50

1875/1875  8s 5ms/step - accuracy: 0.9908 - loss: 0.0277 - val\_accuracy: 0.9788 - val\_loss: 0.0816


Epoch 10/50

1875/1875  11s 5ms/step - accuracy: 0.9911 - loss: 0.0266 - val\_accuracy: 0.9777 - val\_loss: 0.0896


Epoch 11/50

1875/1875  10s 5ms/step - accuracy: 0.9926 - loss: 0.0230 - val\_accuracy: 0.9808 - val\_loss: 0.0795

Epoch 12/50

1875/1875  10s 5ms/step - accuracy: 0.9921 - loss: 0.0238 - val\_accuracy: 0.9812 - val\_loss: 0.0801


Epoch 13/50

1875/1875  9s 5ms/step - accuracy: 0.9925 - loss: 0.0237 - val\_accuracy: 0.9808 - val\_loss: 0.0876


Epoch 14/50

1875/1875  10s 5ms/step - accuracy: 0.9933 - loss: 0.0199 - val\_accuracy: 0.9793 - val\_loss: 0.0890


Epoch 15/50

1875/1875  10s 5ms/step - accuracy: 0.9931 - loss: 0.0199 - val\_accuracy: 0.9803 - val\_loss: 0.0853


Epoch 16/50

1875/1875  10s 5ms/step - accuracy: 0.9933 - loss: 0.0202 - val\_accuracy: 0.9818 - val\_loss: 0.0849

Epoch 17/50


1875/1875  9s 5ms/step - accuracy: 0.9944 - loss: 0.0198 - val\_accuracy: 0.9817 - val\_loss: 0.0842

Epoch 18/50

1875/1875  9s 5ms/step - accuracy: 0.9944 - loss: 0.0165 - val\_accuracy: 0.9815 - val\_loss: 0.0838

curacy: 0.9806 - val\_loss: 0.0949

Epoch 19/50

**1875/1875**  **12s** 5ms/step - accuracy: 0.9929 - loss: 0.0225 - val\_a

ccuracy: 0.9807 - val\_loss: 0.0915

Epoch 20/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9945 - loss: 0.0173 - val\_a


ccuracy: 0.9827 - val\_loss: 0.0871

Epoch 21/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9945 - loss: 0.0172 - val\_a

ccuracy: 0.9826 - val\_loss: 0.0843

Epoch 22/50

**1875/1875**  **9s** 5ms/step - accuracy: 0.9947 - loss: 0.0173 - val\_ac

curacy: 0.9808 - val\_loss: 0.1027

Epoch 23/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9947 - loss: 0.0182 - val\_a

ccuracy: 0.9827 - val\_loss: 0.1024

Epoch 24/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9949 - loss: 0.0160 - val\_a

ccuracy: 0.9806 - val\_loss: 0.1109

Epoch 25/50

**1875/1875**  **11s** 6ms/step - accuracy: 0.9952 - loss: 0.0162 - val\_a

ccuracy: 0.9838 - val\_loss: 0.0914

Epoch 26/50

**1875/1875**  **18s** 4ms/step - accuracy: 0.9953 - loss: 0.0150 - val\_a

ccuracy: 0.9832 - val\_loss: 0.0974

Epoch 27/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9968 - loss: 0.0102 - val\_a


ccuracy: 0.9803 - val\_loss: 0.1123

Epoch 28/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9955 - loss: 0.0157 - val\_a


ccuracy: 0.9824 - val\_loss: 0.1187

Epoch 29/50

**1875/1875**  **9s** 5ms/step - accuracy: 0.9958 - loss: 0.0151 - val\_ac

curacy: 0.9766 - val\_loss: 0.1487

Epoch 30/50

**1875/1875**  **9s** 5ms/step - accuracy: 0.9951 - loss: 0.0185 - val\_ac

curacy: 0.9814 - val\_loss: 0.1170

Epoch 31/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9958 - loss: 0.0148 - val\_a

ccuracy: 0.9823 - val\_loss: 0.1115

Epoch 32/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9957 - loss: 0.0151 - val\_a


ccuracy: 0.9809 - val\_loss: 0.1226

Epoch 33/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9954 - loss: 0.0158 - val\_a

ccuracy: 0.9808 - val\_loss: 0.1267

Epoch 34/50

**1875/1875**  **8s** 4ms/step - accuracy: 0.9965 - loss: 0.0122 - val\_ac

curacy: 0.9823 - val\_loss: 0.1396

Epoch 35/50

**1875/1875**  **11s** 5ms/step - accuracy: 0.9951 - loss: 0.0191 - val\_a

ccuracy: 0.9840 - val\_loss: 0.1247

Epoch 36/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9963 - loss: 0.0144 - val\_a

ccuracy: 0.9823 - val\_loss: 0.1289

Epoch 37/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9956 - loss: 0.0169 - val\_a


ccuracy: 0.9827 - val\_loss: 0.1220

Epoch 38/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9966 - loss: 0.0112 - val\_a

ccuracy: 0.9849 - val\_loss: 0.1187

Epoch 39/50

**1875/1875**  **9s** 5ms/step - accuracy: 0.9963 - loss: 0.0149 - val\_a

curacy: 0.9815 - val\_loss: 0.1373

Epoch 40/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9963 - loss: 0.0115 - val\_a


ccuracy: 0.9839 - val\_loss: 0.1292

Epoch 41/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9971 - loss: 0.0104 - val\_a

ccuracy: 0.9839 - val\_loss: 0.1237

Epoch 42/50

**1875/1875**  **8s** 4ms/step - accuracy: 0.9965 - loss: 0.0116 - val\_a

curacy: 0.9823 - val\_loss: 0.1377

Epoch 43/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9960 - loss: 0.0145 - val\_a

ccuracy: 0.9808 - val\_loss: 0.1488

Epoch 44/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9957 - loss: 0.0149 - val\_a


ccuracy: 0.9806 - val\_loss: 0.1445

Epoch 45/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9966 - loss: 0.0119 - val\_a

ccuracy: 0.9816 - val\_loss: 0.1456

Epoch 46/50

**1875/1875**  **8s** 5ms/step - accuracy: 0.9959 - loss: 0.0153 - val\_a

curacy: 0.9832 - val\_loss: 0.1369

Epoch 47/50

**1875/1875**  **11s** 5ms/step - accuracy: 0.9966 - loss: 0.0116 - val\_a

ccuracy: 0.9823 - val\_loss: 0.1363

Epoch 48/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9964 - loss: 0.0121 - val\_a


ccuracy: 0.9830 - val\_loss: 0.1331

Epoch 49/50

**1875/1875**  **10s** 5ms/step - accuracy: 0.9967 - loss: 0.0144 - val\_a

ccuracy: 0.9817 - val\_loss: 0.1473

Epoch 50/50

**1875/1875**  **9s** 5ms/step - accuracy: 0.9967 - loss: 0.0114 - val\_a

curacy: 0.9807 - val\_loss: 0.1792

After completing the grid search the best value for our hidden size is 200 has it gave us the lowest loss and highest accuracy.